

Are LLMs All We Need?

A "Courtroom" Model for Hybrid AI

Author: John Repsys

With editorial collaboration from: ChatGPT (models GPT-4o and o3, 2025)

Date: April 2025

1. Introduction: The Promise and Limits of Association

In recent years, Large Language Models (LLMs) have displayed remarkable abilities to generate nuanced and context-rich outputs. Their strength lies in association: the capacity to link words, ideas, and concepts based on learned patterns from massive corpora.

Yet association by itself cannot carry the burden of truth. LLMs can suggest, synthesize, and speculate — but without a structured form of verification, they often drift into what we call hallucination. Some critics argue that we might never reach human-like intelligence if we rely on LLMs alone.

This paper proposes that associative intelligence, while powerful, is only one side of reasoning. To move forward, we need systems that integrate:

- The generative freedom of LLMs,
- A persistent, precedent-like memory of established facts (e.g., FAISS/Parquet),
- The constraint of real-world observation and human judgment.

Together, these form a dynamic framework — a *courtroom*, not just a canvas.

2. The Courtroom Model: Legal Precedent & Debate

When humans think, we do not just generate ideas — we challenge them. We dream in association, but we argue with structure. This interplay is core to science, philosophy, and everyday problem-solving.

In this model:

- LLMs serve as the intuition engine, offering possibilities, patterns, and analogies.
- FAISS/Parquet-like systems represent the legal precedent — a persistent, structured memory of facts, embeddings, or prior rulings.
- Real-world observation and human input currently form the courtroom discussion, where conflicting arguments are presented, verified, or refuted — though one of our greatest challenges may be automating parts of this process without losing its rigor.

Once an LLM suggests something novel, it goes before this ‘court’ — a combination of structured memory and open debate (human or multi-agent). This approach helps ensure that not every new association automatically becomes accepted truth.

We can imagine the knowledge frontier as a curve — expanding outward with each novel association, but pulled back when ideas fail to hold. The courtroom model does not just test the new; it re-evaluates the old. In this way, progress is not a straight line, but a dynamic boundary — inching outward where truth endures, folding inward where illusion is exposed. In practice, we are bound by energy constraints — we cannot explore every possibility infinitely, and must prioritize which cases are argued in depth.

3. A Practical Implementation: Financial Data as a Reasoning Playground

While much of the “LLM + formal proof” conversation revolves around pure math (e.g., Lean), the courtroom model has broader applications. One example is financial data, where ambiguity, structure, and semantics collide.

- Line items in SEC filings vary widely in naming and format. LLMs can provide initial opinions on relationships with known metrics.
- Vector embeddings provide semantic similarity checks with known precedent.
- FAISS offers nearest-neighbor searches over high-dimensional vectors.
- Parquet files act as version-aware memory (akin to legal precedent) for validated items.
- Human review introduces friction, preventing illusions from passing as fact.

This architecture becomes a working prototype of the courtroom model itself: association meets structure, guided by friction.

4. Are LLMs a Dead End? Converging Error

Yann LeCun and others have argued that simply scaling LLMs may lead to diminishing returns if they remain purely associative. Indeed, some references imply an error formula, $P(\text{correct}) = (1 - e)^n$, that tends toward 0 rather than converging to stable truths — especially in complex domains.

But that does not mean we must discard what we have built and start over. By integrating a persistent memory (like legal precedent) and real-world constraints (courtroom debate), we can anchor LLM insights and keep them from drifting. The convergence can be tamed if each new claim is forced to verify itself in a structured domain.

Meanwhile, the question “Are LLMs a dead end?” continues as a courtroom discussion of its own. Initially, some speculated that superintelligence was near. Then others, like LeCun, pushed back. The debate remains open — and that is exactly the point: *we need friction to keep exploring ideas*. LLMs deserve their own path of exploration. While some critics argue that LLMs have stalled proper AI advancement, it may be more accurate to say we are finally learning to build one side of the brain. That work should continue — as we learn to build the other side.

5. Multi-LLM Dialogues: An Interim "Court"

Some researchers propose pitting multiple LLMs against each other in a “committee of models.” They debate or critique each other to reduce hallucinations. While promising, such committees can still converge on collective illusions, especially if they share the same biases or training distributions.

That is where external friction (real-world data, symbolic checks, or human oversight) remains crucial. We need a memory that cannot be overturned by the illusions that created it. In legal terms, you cannot rewrite established precedent without a higher burden of proof.

6. Architectural Analogues

A notable example of the Courtroom Model in action is AlphaGo, developed by DeepMind. While not explicitly framed this way at the time, AlphaGo's architecture embodies the same tripartite pattern:

- A policy network generates likely moves (analogous to LLM-style associative reasoning).
- A value network evaluates the resulting position (structured memory or precedent).
- Self-play provides the iterative, adversarial challenge (analogous to courtroom-style filtering and argument)

Despite being a fully automated system, AlphaGo reflects the same structural rhythm: generate, evaluate, and challenge. This convergence supports the idea that the Courtroom Model is not prescriptive, but descriptive — a lens for recognizing effective reasoning architectures wherever they emerge. It offers a retrospective framework to describe such structures — and perhaps guide the design of future systems that balance creativity with constraint.

Another related effort is the combination of Lean + LLMs in formal mathematics. There, the LLM suggests steps or statements, while the Lean proof assistant rigorously checks their correctness — mimicking the generation/challenge dichotomy at a symbolic level. Both AlphaGo and Lean + LLM workflows demonstrate that progress often arises from structured friction, not just creative leaps.

7. Open Questions

1. **Is abstraction a distinct step in hybrid reasoning?** When a model or human finds no exact precedent, it must sometimes create a new one — compressing multiple associations into a reusable concept. Is this a unique layer in the reasoning process? Should pattern recognition and abstraction be formal components in systems like this one?
2. **Can LLMs + real-world observation produce entirely new theories?** If so, how do we detect genuine novelty, and how do we measure validity?
3. **Can the Courtroom Model generalize to other real-world domains and AI integrations?** Healthcare, policy, engineering — any area that needs scale, creativity, and robust checks. Could this framework point toward a general pattern for hybrid AI?
4. **Other than using LLMs en masse, how can we automate analogous versions of Lean in areas other than mathematics?** The group of LLMs approach is a good first try, but it lacks the precision many fields demand for automation to succeed.

5. **How do we maintain the tension between speed and scrutiny?** Associative models want to generate rapidly; symbolic structures demand thorough verification. Speed wins until it fails. Scrutiny survives because it must.
6. **What does success look like in a system that's always evolving?** As new precedents form, do we prune old ones? How do we adapt long-standing truths? The answer may lie in how far our energy limitations let us revisit and revise what we already hold to be true.

8. Toward Hybrid Intelligence

If we want machines that do not just speak, but reason, we must build:

- Systems that suggest yet also remember
- Systems that associate yet also verify
- Systems that propose yet also fail with friction

This is not just a technical project; it is a philosophical one. It mirrors humanity's own process of knowledge-building: sometimes we push boundaries, sometimes we retreat. Progress emerges from association, tested and refined by structured debate.

9. Conclusion: A New Frontier

Pure association is not enough — but it is not a dead end. Symbolic structure alone is not sufficient — but it is an essential grounding force. Real-world observation provides the friction, cost, and stakes that compel ideas to prove themselves.

Whether we apply this to financial data, formal math, or any domain where illusions are costly, the principle remains: Good reasoning is associative, structured, and tested by reality.

We do not need to throw out the progress LLMs have made. *We just need to give them a courtroom.*

References

- Vaswani et al., *Attention Is All You Need*, 2017.
- LeCun, Y. *Introduction to Latent Variable Energy-Based Models: A Path Towards Autonomous Machine Intelligence*, 2023.
- Buzzard, K. Lectures on formal proof and Lean (YouTube).
- Marcus, G. *The Next Decade in AI: Four Steps Toward Robust Artificial Intelligence*, 2020.

Acknowledgments

- Inspired by structured data design and the vision of symbolic + semantic harmony. Special thanks to ongoing conversations about Lean, neuro-symbolic AI, multi-agent dialogue, and the interplay of association vs. verification.

