

# **Are LLMs All We Need?**

## **A "Courtroom" Model for Hybrid AI**

**Author:** John Repsys

**With editorial collaboration from:** ChatGPT (models GPT-4o and o3, 2025)

**Date:** April 2025

# 1. Introduction: The Promise and Limits of Association

Large Language Models (LLMs) have demonstrated remarkable abilities to generate context-rich outputs through association, linking words and concepts based on patterns learned from large corpora. However, this power comes with limitations. While association can drive creativity, it does not guarantee the truth. LLMs often generate plausible-sounding but incorrect conclusions—a phenomenon known as hallucination. Without a structured verification system, these models risk perpetuating errors and amplifying inaccuracies that may seem credible at first glance, undermining their reliability.

This paper proposes that LLMs are powerful tools, but they must be paired with systems that integrate:

- The generative freedom of LLMs,
- A persistent, precedent-like memory of established facts (e.g., FAISS/Parquet),
- The dynamic validation of generated outputs, whether through external feedback (e.g., human judgment or real-world data) or internal checks (e.g., iterative self-play in closed domains).

Together, these form a dynamic framework — a courtroom, not just a canvas. Unlike systems like ReConcile or MA-LoT, which utilize multi-agent approaches or predefined rules to aggregate insights, the Courtroom Model emphasizes a perpetual validation loop that brings external feedback, evolving precedent, and iterative re-evaluation into the equation. This integration ensures that LLMs' ability to generate novel associations is balanced by rigorous verification—actively preventing the propagation of errors before they can take root, and ensuring that only those ideas that stand up to scrutiny are accepted.

The Courtroom Model is not based on fundamentally new components but rather on established technologies like LLMs, persistent memory systems such as FAISS/Parquet, and human oversight. What differentiates this model is its meta-level abstraction, which reconfigures these elements into a dynamic system of ongoing validation and refinement. Just as mathematicians generalize concepts into more universal structures, the Courtroom Model transforms how we approach AI reasoning by continuously testing, challenging, and refining knowledge—not just generating it.

## 2. The Courtroom Model: Legal Precedent & Debate

Humans think by generating ideas and challenging them. In scientific and philosophical inquiry, ideas are not just created—they are debated, tested, and refined. As discussed earlier, the Courtroom Model integrates LLMs for idea generation, FAISS/Parquet for memory, and feedback loops for validation. In this section, we'll explore how these elements interact in greater detail.

Once an LLM generates an idea, it enters the 'courtroom'—where it is tested against structured precedent and subjected to rigorous debate. Ideas are challenged and refined, ensuring that only those that withstand scrutiny are accepted as truth. This approach ensures that new associations are not blindly accepted but are examined and tested continually, resulting in dynamic, evolving knowledge. In this way, the Courtroom Model doesn't merely reinforce existing knowledge—it ensures conclusions are constantly scrutinized and continuously tested, preventing them from being accepted without

question. This re-evaluation ensures that knowledge is constantly refined and that assumptions are actively challenged in light of new evidence.

We can imagine the knowledge frontier as a quarter circle in the first quadrant, with the curve representing the boundary of our current understanding. The area inside the curve symbolizes the cumulative knowledge at any given time. As new associations or innovations emerge, the frontier may expand outward, pushing the curve further along—provided these new ideas meet verification checks. However, the system doesn't just push the frontier outward; it continually revisits and checks the knowledge inside the curve for potential errors or hallucinations. In this way, progress is not a linear trajectory but a dynamic boundary that expands outward where truth endures, while folding inward when inaccuracies are identified. As we explore this frontier, we are mindful of energy constraints, prioritizing which areas require deeper scrutiny and ensuring that exploration remains focused and efficient.

The Courtroom Model introduces a hybrid framework that distinguishes itself from current AI methodologies by integrating associative reasoning (LLMs) with the persistent structure of legal precedents (FAISS/Parquet) and dynamic validation. This combination enables the model to adapt, incorporate human oversight, and apply continuous scrutiny—features largely absent in systems like ReConcile or MA-LoT.

Rather than requiring manual review of every idea, the Courtroom Model automatically flags associations that meet certain thresholds of uncertainty or inconsistency with established facts. This approach allows humans to focus on the most critical evaluations, ensuring that their intervention is directed toward areas where it adds the most value.

### **3. A Practical Implementation: Financial Data as a Reasoning Playground**

While much of the LLM + formal proof conversation focuses on pure math (e.g., Lean), the Courtroom Model has broad applications, including complex domains like financial data. In this space, ambiguity and structure meet, and semantic checks become crucial:

- SEC filings often contain inconsistent naming and formats. LLMs can provide initial opinions on relationships between known metrics.
- Vector embeddings allow semantic similarity checks with prior precedent, using metrics such as L2-distance or cosine similarity to quantify the closeness between new ideas and established knowledge. These checks focus on identifying close matches to validated concepts and flagging potential outliers, ensuring that the system not only finds similarities but also highlights discrepancies that might indicate the need for deeper scrutiny or re-evaluation.
- FAISS performs nearest-neighbor searches over high-dimensional vectors, helping to identify similar concepts based on predefined thresholds.
- Parquet files serve as version-aware memory, storing validated items, like legal precedent.

- In financial analysis, when a generative model suggests a correlation or metric, human analysts review it against historical data, verify the logic, and confirm its validity. This friction ensures errors do not become accepted facts.

This process embodies the Courtroom Model: association meets structure, refined and verified by friction. One of the key innovations of the Courtroom Model is its emphasis on the continuous re-evaluation of previously accepted conclusions. Unlike systems such as ReConcile and MA-LoT, which may stabilize and converge over time, the Courtroom Model actively revisits past conclusions, ensuring that all knowledge is up-to-date and consistent with the latest data. This perpetual loop of validation ensures that knowledge is not only updated in response to new data but is also actively managed by prioritizing certain claims for re-evaluation, especially as new evidence challenges older beliefs. The Courtroom Model actively manages which ideas undergo deeper scrutiny, using real-world data, symbolic checks, human oversight, and similarity metrics like L2-distances or cosine similarity to determine when and how older claims need to be revisited.

In practice, the Courtroom Model could leverage dynamic prioritization algorithms to determine which associations or precedents are most in need of re-evaluation. This approach would enable the model to scale effectively to large datasets, focusing scrutiny on the most impactful claims and reducing computational load while maintaining rigorous validation. While this is not yet implemented, it represents a key area for future development, enhancing the system’s efficiency and scalability.

## 4. Are LLMs a Dead End? Converging Error

Critics like Yann LeCun argue that LLMs, if scaled purely based on association, may experience diminishing returns, especially in complex domains. Indeed, some references imply an error formula,  $P(\text{correct}) = (1 - e)^n$ , that tends toward 0 rather than converging to stable truths. LeCun also notes the inability of LLMs to represent continuous high-dimensional spaces that characterize the real world, potentially limiting their ability to deal with complex, nonlinear data.

However, this does not mean we must abandon LLMs. Rather than expecting LLMs to handle everything on their own, the Courtroom Model allows them to excel at generating ideas and associations, while supplementing them with additional systems that anchor their outputs and guide them toward more reliable outcomes. By integrating persistent memory (like legal precedent) and real-world constraints (courtroom debate), the model offers a framework for continuously verifying and refining LLM-generated associations, ensuring they align with established facts and real-world data. This helps overcome challenges posed by complex, high-dimensional data, and prevents errors from propagating.

Meanwhile, the question “Are LLMs a dead end?” continues to evolve as part of the broader debate on AI’s future. Some initially speculated that superintelligence was near, while others, like LeCun, pushed back, suggesting that LLMs’ potential was overstated. The debate remains open—and that is exactly the point: we need friction to continue exploring ideas and improving models. The Courtroom Model provides a way forward by leveraging LLMs in combination with other systems, ensuring that we don’t discard their value but integrate them into a broader framework that allows them to evolve

meaningfully. While critics argue that LLMs have stalled AI progress, it may be more accurate to say that we're just learning to build one side of the brain, and as we expand our approach, we must continue building the other side as well.

## 5. Multi-LLM Dialogues: An Interim "Court"

Some researchers propose pitting multiple LLMs against each other in a committee of models to reach a consensus. While this approach can help reduce individual model biases, it still carries the risk of collective hallucination—even when the models are trained on different datasets or don't share the same biases. Because LLMs generate associative reasoning based on patterns, they can unintentionally converge on similar, plausible-sounding errors, even without shared flaws in their training data. This collective error is a form of hallucination that can persist if left unchecked. This is where the Courtroom Model offers a critical advantage: it integrates both internal validation and external friction. In this way, the Courtroom Model prevents multi-LLM systems from reinforcing errors or hallucinating collectively, ensuring that new ideas undergo rigorous validation before being accepted. This adds an extra layer of scrutiny that multi-LLM systems lack, resulting in more accurate and reliable outputs.

## 6. Architectural Analogues

The Courtroom Model is not just a theoretical framework—it has analogues in existing AI systems. For example, AlphaGo by DeepMind used a similar process of generating moves (LLM-style reasoning) and evaluating outcomes (structured memory/precedent). While AlphaGo operates in a closed domain, relying on self-play for evaluation, the Courtroom Model extends this concept by incorporating both internal validation (within systems like FAISS/Parquet) and external checks (such as human judgment or real-world data) when needed. This flexibility makes the Courtroom Model applicable not just to closed systems like AlphaGo, but to open systems that require dynamic and evolving feedback to ensure accuracy. While not explicitly framed this way at the time, AlphaGo's architecture reflects the same tripartite pattern:

- A policy network generates likely moves (analogous to LLM-style associative reasoning).
- A value network evaluates the resulting position (structured memory or precedent).
- Self-play provides the iterative, adversarial challenge (analogous to courtroom-style filtering and argument)

Despite being a fully automated system, AlphaGo reflects the same structural rhythm: generate, evaluate, and challenge. This convergence supports the idea that the Courtroom Model is not prescriptive, but descriptive — a lens for recognizing effective reasoning architectures wherever they emerge. It offers a retrospective framework to describe such structures — and perhaps guide the design of future systems that balance creativity with constraint.

Another related effort is the combination of Lean + LLMs in formal mathematics. In this setup, the LLM generates proposals for mathematical steps or theorems, which are then checked for logical consistency by the Lean proof assistant. Lean ensures these steps adhere to formal logic before they are

accepted, while the Courtroom Model similarly ensures LLM-generated ideas are rigorously validated by structured processes and real-world feedback. This mirrors the generation/challenge dichotomy in the Courtroom Model, where the LLM suggests associations or ideas, but they are tested against pre-existing knowledge (precedent) and scrutinized for consistency with real-world data. Just as Lean ensures that mathematical steps are sound before they are accepted as valid, the Courtroom Model ensures that LLM-generated ideas are continuously tested, refined, and validated by structured processes and external feedback.

Both AlphaGo and Lean + LLM workflows demonstrate that progress often arises from structured friction, not just creative leaps. In the same way, the Courtroom Model allows for creativity (via LLMs) to drive idea generation, but that creativity is continually tested and refined through structured validation and external challenge, ensuring only robust, verified ideas become part of the model's evolving knowledge.

## 7. Open Questions

1. **Is abstraction a distinct step in hybrid reasoning?** When no exact precedent exists, does the model create a new one by compressing multiple associations? Should pattern recognition and abstraction be formal components in systems like this one?
2. **Can LLMs + real-world observation produce entirely new theories?** If so, how do we detect genuine novelty, and how do we measure validity?
3. **Can the Courtroom Model generalize to other real-world domains and AI integrations?** Healthcare, policy, engineering — any area that needs scale, creativity, and robust checks. Could this framework point toward a general pattern for hybrid AI?
4. **Other than using large groups of LLMs, how can we automate analogous versions of Lean in areas other than mathematics?** The group of LLMs approach is a good first try, but it lacks the precision many fields demand for automation to succeed.
5. **How do we maintain the tension between speed and scrutiny?** Associative models want to generate rapidly; symbolic structures demand thorough verification. Speed wins until it fails. Scrutiny survives because it must.
6. **What does success look like in a system that's always evolving?** As new precedents form, do we prune old ones? How do we adapt long-standing truths? The answer may lie in how far our energy limitations let us revisit and revise what we already hold to be true.

## 8. Toward Hybrid Intelligence

To achieve reasoning, machines must:

- Suggest and remember
- Associate and verify
- Propose and fail with friction

This is not just a technical project; it is a philosophical one. It mirrors humanity’s own process of knowledge-building: sometimes we push boundaries, sometimes we retreat. Progress emerges from association, tested and refined by structured debate.

## 9. Conclusion: A New Frontier

Pure association is not enough — but it is not a dead end. Symbolic structure alone is not sufficient — but it is an essential grounding force. Real-world observation provides the friction, cost, and stakes that compel ideas to prove themselves.

Whether we apply this to financial data, formal math, or any domain where illusions are costly, the principle remains: Good reasoning is associative, structured, and tested by reality.

To ensure that AI can reliably and consistently contribute to our decision-making, we must move beyond mere association and toward systems that integrate creativity with rigorous validation. The Courtroom Model offers a critical path forward, empowering AI to not only generate ideas but to validate, challenge, and refine them in a continuous loop. This dynamic approach is essential for building AI that is not just intelligent, but also trustworthy—ensuring that AI evolves in a way that can be reliably trusted across various real-world applications.

However, the Courtroom Model is just a step forward in AI reasoning. It represents a framework that combines generative power with structured validation. As the field evolves, it serves as a foundation upon which we can build more robust systems, continually refining and improving AI to meet the complex demands of the real world.

## References

- Vaswani et al., *Attention Is All You Need*, 2017.
- LeCun, Y. *Introduction to Latent Variable Energy-Based Models: A Path Towards Autonomous Machine Intelligence*, 2023.
- Buzzard, K. Lectures on formal proof and Lean (YouTube).
- Marcus, G. *The Next Decade in AI: Four Steps Toward Robust Artificial Intelligence*, 2020.

## Acknowledgments

- Inspired by structured data design and the vision of symbolic + semantic harmony. Special thanks to ongoing conversations about Lean, neuro-symbolic AI, multi-agent dialogue, and the interplay of association vs. verification.